

Lecture 9: Supervised Learning – Neural Networks

Reading Assignment: EoSL Chapter 11.1, 11.3-11.8

Neural Networks in R Tutorial: <https://www.datacamp.com/community/tutorials/neural-network-models-r>

- Recall dimension reduction methods via derived input methods from Lecture 3: produce *linear* combinations Z_m [$m = 1, \dots, M$] of original inputs X_j to use in place of X_j in the regression.

$$Z_m = \sum_{j=1}^p \phi_{jm} X_j$$

for some constants $\phi_{1m}, \dots, \phi_{pm}$, $m = 1, \dots, M$ where the linear regression model becomes

$$y_i = \beta_0 + \sum_{m=1}^M \beta_m \sum_{j=1}^p \phi_{jm} x_{ij} + \epsilon_i = \beta_0 + \sum_{j=1}^p \beta_j^* x_{ij} + \epsilon_i$$

thus modeling the target as a linear function of the derived features.


- Also recall nonlinear regression methods from Lecture 4 that relax linearity through linear basis expansion in X :

$$f(X) = \sum_{m=1}^M \beta_m h_m(X).$$

- Neural networks yield a large class of models that expand on these ideas.
- Neural networks extract linear combinations of inputs as derived features, then model the target as a *nonlinear* function of these features.
- They are just nonlinear statistical models – despite their surrounding hype as magical and mysterious!
- A neural network is just a two-stage regression or classification model that is typically represented by a network diagram.
- “Neural network” comes from being first developed as models for the human brain – each unit represents a neuron and the connections are the synapses.
- Lets start with a foundational concept/construct in (artificial) neural networks.

* **Perceptrons**

- A perceptron is an artificial neuron developed in the field of AI.
- Perceptrons take binary inputs X_1, \dots, X_p to produce a binary output.
- *Example:* $p = 3$.




- A simple rule to compute the neurons 0-1 output:

$$\text{output} = \begin{cases} 0 & \text{if } \sum_{j=1}^p w_j x_j \leq t \\ 1 & \text{if } \sum_{j=1}^p w_j x_j > t, \end{cases}$$

where w_1, \dots, w_p are weights (reflecting cost/importance of an input on the decision) and t is a real number threshold (the parameter of the neuron).

- The perceptron is a device to make binary decisions by weighting evidence; a model for decision-making.
- It is natural to envision a network of perceptrons.



- The column of perceptrons – called the first layer – is making simple decisions by weighting the input layer.
- *Note:* nodes send signals in only one direction thus called a feed forward network. Recurrent neural networks allow loops in the network.

- A simpler way to write the simple rule

$$\text{output} = \begin{cases} 0 & \text{if } w_j^T x_j + b \leq 0 \\ 1 & \text{if } w_j^T x_j + b > 0, \end{cases}$$

where b is the perceptron's bias that measures how easy it is to get the perceptron to output a 1.

- Use learning algorithms to tune the weights and the bias of a network of artificial neurons.
- Lets use this to formulate the plain vanilla neural net as a learning algorithm.

* **Single Hidden Layer (Single Layer Perceptron) Network**

- A single hidden layer neural network diagram for regression (typically Y_1 only) or classification (Y_1, \dots, Y_K).




FIGURE 11.2. Schematic of a single hidden layer, feed-forward neural network.

- All neural networks have an input and output layer.
- Edges “link” the nodes – a “fully connected” network ensures that all nodes in the lower layer contribute to each layer above.
- Derived features, a.k.a. perceptrons or hidden units, Z_m define the hidden layer. They are created from linear combinations of the inputs

$$Z_m = \sigma(\alpha_{0m} + \alpha_m^T X) \quad (1)$$

for $m = 1, \dots, M$ [from earlier notation $\alpha_j = w_j$ and $\alpha_0 = b$].

- In original formulations σ is a step function – i.e. neurons fire when the signal exceeded a threshold – but it is more useful for statistical modeling (and optimization) to use a smooth threshold function.
- That is, for the perceptron a small change in input values can cause a large change in output; a better solution would be to output a continuum of values.
- The activation function $\sigma(\nu)$ is usually chosen to be the sigmoid (a.k.a. logit, softmax) function

$$\sigma(\nu) = \frac{1}{1 + e^{-\nu}}, \quad (2)$$

a smoothed version of a step function.




FIGURE 11.3. Plot of the sigmoid function $\sigma(v) = 1/(1+\exp(-v))$ (red curve),

- The sigmoid neuron model allows inputs to take any value between 0 and 1 and yields an output that is not just 0 or 1.
- For particularly positive or negative ν , results will be nearly the same as the perceptron (0 or 1); but intermediate values will be spread out.
- The target Y_k is modeled as a function of linear combinations of Z_m

$$f_k(X) = g_k\left(\beta_{0k} + \beta_k^T Z\right) \quad (3)$$

for $k = 1, \dots, K$ where $Z = (Z_1, \dots, Z_m)$.

- The output function $g_k(T_k)$ allows a final transformation of the outputs $T_k = \beta_{0m} + \beta_k^T Z$. Typically chosen to be

▷ identity function for regression: $g_k(T_k) = T_k$,

▷ softmax/multilogit function for classification: $g_k(T_k) = \frac{e^{T_k}}{\sum_{\ell=1}^K e^{T_\ell}}$.

- In summary, components of neural networks:
 - ▷ *input layer*: p features X_1, \dots, X_p ,
 - ▷ *hidden layer* (1): $Z_m = \sigma(\alpha_{0m} + \alpha_m^T X)$ for $m = 1, \dots, M$,
 - ▷ *output layer* (3): $f_k(X) = g_k(\beta_{0k} + \beta_k^T Z) = g_k(T_k)$,
 - ▷ *activation function* (2): $\sigma(\alpha_{0m} + \alpha_m^T X) = \frac{1}{1 + e^{-(\alpha_{0m} + \alpha_m^T X)}}$, and
 - ▷ *link function* (3): g_k , e.g. logit or identity.
- Considering Z_m as a pre-defined basis expansion of original inputs and $\sigma(\nu)$ assumed to be the identity function \Rightarrow standard nonlinear models from Lecture 4, i.e. nonlinear generalization of the linear model.
- But now introducing the nonlinear transformation σ greatly expands the class of models.
- Also in neural networks, *the basis functions are learned from the data* jointly with the target model; the Z_m are not directly observed – they are hidden!

* **Fitting Neural Networks (Single Layer)**

- Need to estimate the unknown parameters, a.k.a. weights, denoted by θ which consists of

$$\begin{aligned} \{\alpha_{0m}, \alpha_m; m = 1, \dots, M\} &\rightarrow M(p + 1) \text{ weights,} \\ \{\beta_{0k}, \beta_k; k = 1, \dots, K\} &\rightarrow K(M + 1) \text{ weights.} \end{aligned}$$

where, recall, M is the number of linear combinations of X , and K is the dimension of the target.

- As always, we are looking to minimize error (loss)

$$R(\theta) = \sum_{k=1}^K \sum_{i=1}^N L(y_{ik}, f_k(x_i)).$$

- *Example*: classification using cross-entropy/deviance loss

$$R(\theta) = - \sum_{i=1}^N \sum_{k=1}^K y_{ik} \log f_k(x_i),$$

with the corresponding classifier $G(x) = \max_k f_k(x)$. *Note*: softmax activation function and deviance error \Rightarrow logistic regression model in Z so can use ML estimation.

- A generic approach to minimizing $R(\theta)$ is by gradient descent, a.k.a. back-propagation in this setting.
- This approach is reasonable because the form of the gradient can easily be derived via chain rule – the model is defined in a compositional form.
- *Example*: regression using squared error loss for multivariate target $Y =$

$$(Y_1, \dots, Y_K),$$

$$\begin{aligned} R(\theta) &= \sum_{i=1}^N R_i = \sum_{i=1}^N \sum_{k=1}^K \left(y_{ik} - f_k(x_i) \right)^2, \\ &= \sum_{i=1}^N \sum_{k=1}^K \left(y_{ik} - g_k(\beta_{0k} + \beta_k^T z_i) \right)^2, \\ &= \sum_{i=1}^N \sum_{k=1}^K \left(y_{ik} - g_k \left(\beta_{0k} + \sum_{m=1}^M \beta_{km} z_{im} \right) \right)^2, \\ &= \sum_{i=1}^N \sum_{k=1}^K \left(y_{ik} - g_k \left(\beta_{0k} + \sum_{m=1}^M \beta_{km} \sigma(\alpha_{0m} + \alpha_m^T x_i) \right) \right)^2, \end{aligned}$$

yielding derivatives (indicating how quickly loss changes with the weights)

$$\begin{aligned} \frac{\partial R_i}{\partial \beta_{km}} &= \frac{\partial R_i(\theta)}{\partial g_k} \frac{\partial g_k}{\partial T_k} \frac{\partial T_k}{\partial \beta_{km}} \\ &= \left[-2 (y_{ik} - f_k(x_i)) g'_k(\beta_{0k} + \beta_k^T z_i) \right] z_{im} \\ &= \delta_{ki} z_{im}, \\ \frac{\partial R_i}{\partial \alpha_{mj}} &= \sum_{k=1}^K \frac{\partial R_i(\theta)}{\partial g_k} \frac{\partial g_k}{\partial T_k} \frac{\partial T_k}{\partial \sigma} \frac{\partial \sigma}{\partial \alpha_{mj}} \\ &= \left[-2 \sum_{k=1}^K (y_{ik} - f_k(x_i)) g'_k(\beta_{0k} + \beta_k^T z_i) \beta_{km} \sigma'(\alpha_{0m} + \alpha_m^T x_i) \right] x_{ij} \\ &= s_{mi} x_{ij}, \end{aligned}$$

so that the gradient descent update at the $(r + 1)$ st iteration is

$$\begin{aligned}\beta_{km}^{(r+1)} &= \beta_{km}^{(r)} - \gamma_r \sum_{i=1}^N \frac{\partial R_i}{\partial \beta_{km}^{(r)}} \\ \alpha_{mj}^{(r+1)} &= \alpha_{mj}^{(r)} - \gamma_r \sum_{i=1}^N \frac{\partial R_i}{\partial \alpha_{mj}^{(r)}}\end{aligned}$$

with learning rate γ_r .

The quantities δ_{ki} and s_{mi} are the “errors” from the current model at the output and hidden layers, respectively and satisfy

$$\begin{aligned}s_{mi} &= -2 \sum_{k=1}^K (y_{ik} - f_k(x_i)) g'_k(\beta_{0k} + \beta_k^T z_i) \beta_{km} \sigma'(\alpha_{0m} + \alpha_m^T x_i) \\ &= \sum_{k=1}^K \left[-2 (y_{ik} - f_k(x_i)) g'_k(\beta_{0k} + \beta_k^T z_i) \right] \beta_{km} \sigma'(\alpha_{0m} + \alpha_m^T x_i) \\ &= \sigma'(\alpha_{0m} + \alpha_m^T x_i) \sum_{k=1}^K \delta_{ki} \beta_{km}\end{aligned}$$

which are called the back-propagation equations.


The updates using gradient descent can be implemented in a *two-pass algorithm* called back-propagation:

- ▷ forward pass: fix weights θ at current value $\hat{\theta}^{(r)}$ and compute $\hat{f}_k(x_i)$,
- ▷ backward pass: fix Z_m, T_k and $f_k(x_i)$ (output) and update θ using a

gradient descent step. Specifically, compute errors δ_{ki} to obtain s_{mi} [i.e. back-propagate], then use errors to compute the gradients to get $\hat{\theta}^{(r+1)}$,

i.e. the back propagation algorithm is a numerical way to compute the gradient to use in the gradient descent numerical optimization algorithm.

- Computational components for deviance loss have same form as for squared error.
- *Intuition*: assess the output error at each step and update the weights of the neural network to reduce error.



- *Advantages*:

- ▷ hidden unit passes/receives information only to/from units with a connection – can parallelize, and

- ▷ can be done as batch learning by summing over all training cases; or online learning processing each observation at a time (cycling through the entire training dataset a specified number of times called epochs), updating gradient after each training case to allow very large training sets.
- Training neural networks is an art – they are generally overparameterized, are a nonconvex optimization problem and can be very unstable.
- There are some very important practical considerations that can make your life easier.

* **Practical Issues in Training Neural Networks**

(1) *Starting Values*

- If the weights are near zero, then the operable part of the sigmoid is roughly linear \Rightarrow network collapses to an approximately linear model.
- Good practice to start model out nearly linear and let the network become nonlinear as the weights increase.
- Starting at exactly zero leads to zero derivatives – algorithm never moves.
- Starting with large weights often leads to poor solutions.

(2) *Overfitting*

- Typically the global minimizer of $R(\theta)$ will overfit [lots of weights!], so

need to add regularization.

- Early developments used an early stopping rule to stop well before approaching the global minimum. Weights start at a highly regularized (linear) solution so this effectively shrinks toward a linear model.
- Can explicitly regularize using weight decay which is analogous to ridge regression for linear models to minimize

$$R(\theta) + \lambda \left(\sum_{k=1}^K \sum_{m=1}^M \beta_{km}^2 + \sum_{m=1}^M \sum_{j=1}^p \alpha_{mj}^2 \right)$$

where $\lambda \geq 0$ is a tuning parameter.

- ↗ $\lambda \Rightarrow$ shrink weights to zero.
- Typically use CV to select λ .
- Other forms of penalty have been proposed.

(3) *Scaling of Inputs*

- Scaling of inputs determines the effective scaling of weights in the input layer and can have a large effect on quality of the final solution.
- Always standardize inputs to have mean zero and standard deviation one.
- Ensures all inputs are treated equally in the regularization process.
- Also can then choose meaningful range for random starting weights: typ-

ical to take random uniform weights over the range $[-0.7, +0.7]$.


(4) *Number of Hidden Units*

- Generally better to have too many hidden units than too few (not enough flexibility).
- With too many, the extraneous can be shrunk to zero with regularization.
- Best to choose reasonable large M and train them with regularization.
- Typical choice is in range of 5 to 100, with the number increasing with p and N .
- Can use CV, but not necessary if you use CV to estimate the regularization parameter λ .

(4) *Multiple Minima*

- The error function $R(\theta)$ is nonconvex with many local minima.
- Final solution is quite dependent on starting weights.
- Best to try many random starting value configurations and choose solution with lowest error.
- Or could use averaging of predictions over the collections of networks or bagging (averages of bootstrapped training data).

- More complicated neural networks can be created with more hidden layers, i.e. multilayer perceptron (MLP).
- Additional layers make simple decisions by weighting the results from the previous layer; they make decisions at a more complex and abstract level than previous layer.



For single layer only [not estimated with gradient descent] with regularization:

— **nnet()** function in **nnet** package in **R** —


Allows multiple layers (vector for **hidden** option) and a variety of algorithms but not regularization:

— **neuralnet()** function in **neuralnet** package in **R** —

Allows multiple layers (vector for **hidden.layers** option) with regularization options and batch options:

— **neuralnetwork()** function in **ANN2** package in **R** —

- What makes neural networks so different and great?
- Scale! As larger neural networks are constructed and trained with more and more data, their performance continues to increase, whereas for other ML methods performance plateaus.



- Computing power continues to improve – able to fit larger and larger neural networks.
- This is the crux of “deep learning” – large neural networks, e.g. MLP – where the graph is deep with many layers.
- But they require more computation power and more data.
- Also require more development time – not “off-the-shelf” (i.e. fast, easy to train). A variety of hyperparameters can be tuned:

- ▷ *neural network structure*: λ , number of layers, M , $\sigma(\cdot)$, $g_k(\cdot)$;
- ▷ *optimization/training algorithm*: loss function, γ_r , starting values, batch size, number of epochs.